

Problem Statement: Cricket Expert Chatbot (with LangSmith)

Build a Cricket Expert Chatbot using LangChain and a Retrieval-Augmented Generation (RAG) pipeline to answer questions about cricket players from a PDF knowledge base.

Requirements

- PDF Ingestion:
 - Load cricket_players.pdf using PyPDFLoader.
 - Split using CharacterTextSplitter (chunk_size=100, chunk_overlap=10).
- Vector Store (choose one):
 - Chroma or FAISS with OllamaEmbeddings (granite-embedding:latest).
 - Store under collection/index: cricket_players.
- LLM Integration (choose one):
 - HuggingFaceEndpoint (e.g., Menlo/Jan-nano-128k) or
 - Local Ollama model (e.g., llama3, mistral).
- Chat Logic:
 - Use RetrievalQA with prompt template.
 - Track conversation with InMemoryChatMessageHistory.
- LangSmith:
 - Enable tracing with LangSmith.
 - Use env vars for API key and project name.

- UI:
 - Gradio interface with:
 - Text input, response area
 - Submit and Clear History buttons
 - Greeting: "Hello, I am CricketBot, your cricket expert. How can I help you today?"

Deliverable

Submit a single `cricket_chatbot.py` that:

- Loads and splits the PDF
- Initializes Chroma or FAISS
- Connects to Ollama or Hugging Face model
- Implements LangSmith tracing
- Runs Gradio UI with history support